



Arhitecturi Paralele Introducere GPU


Lect. Dr. Ing. Cristian Chilipirea – cristian.chilipirea@mta.ro







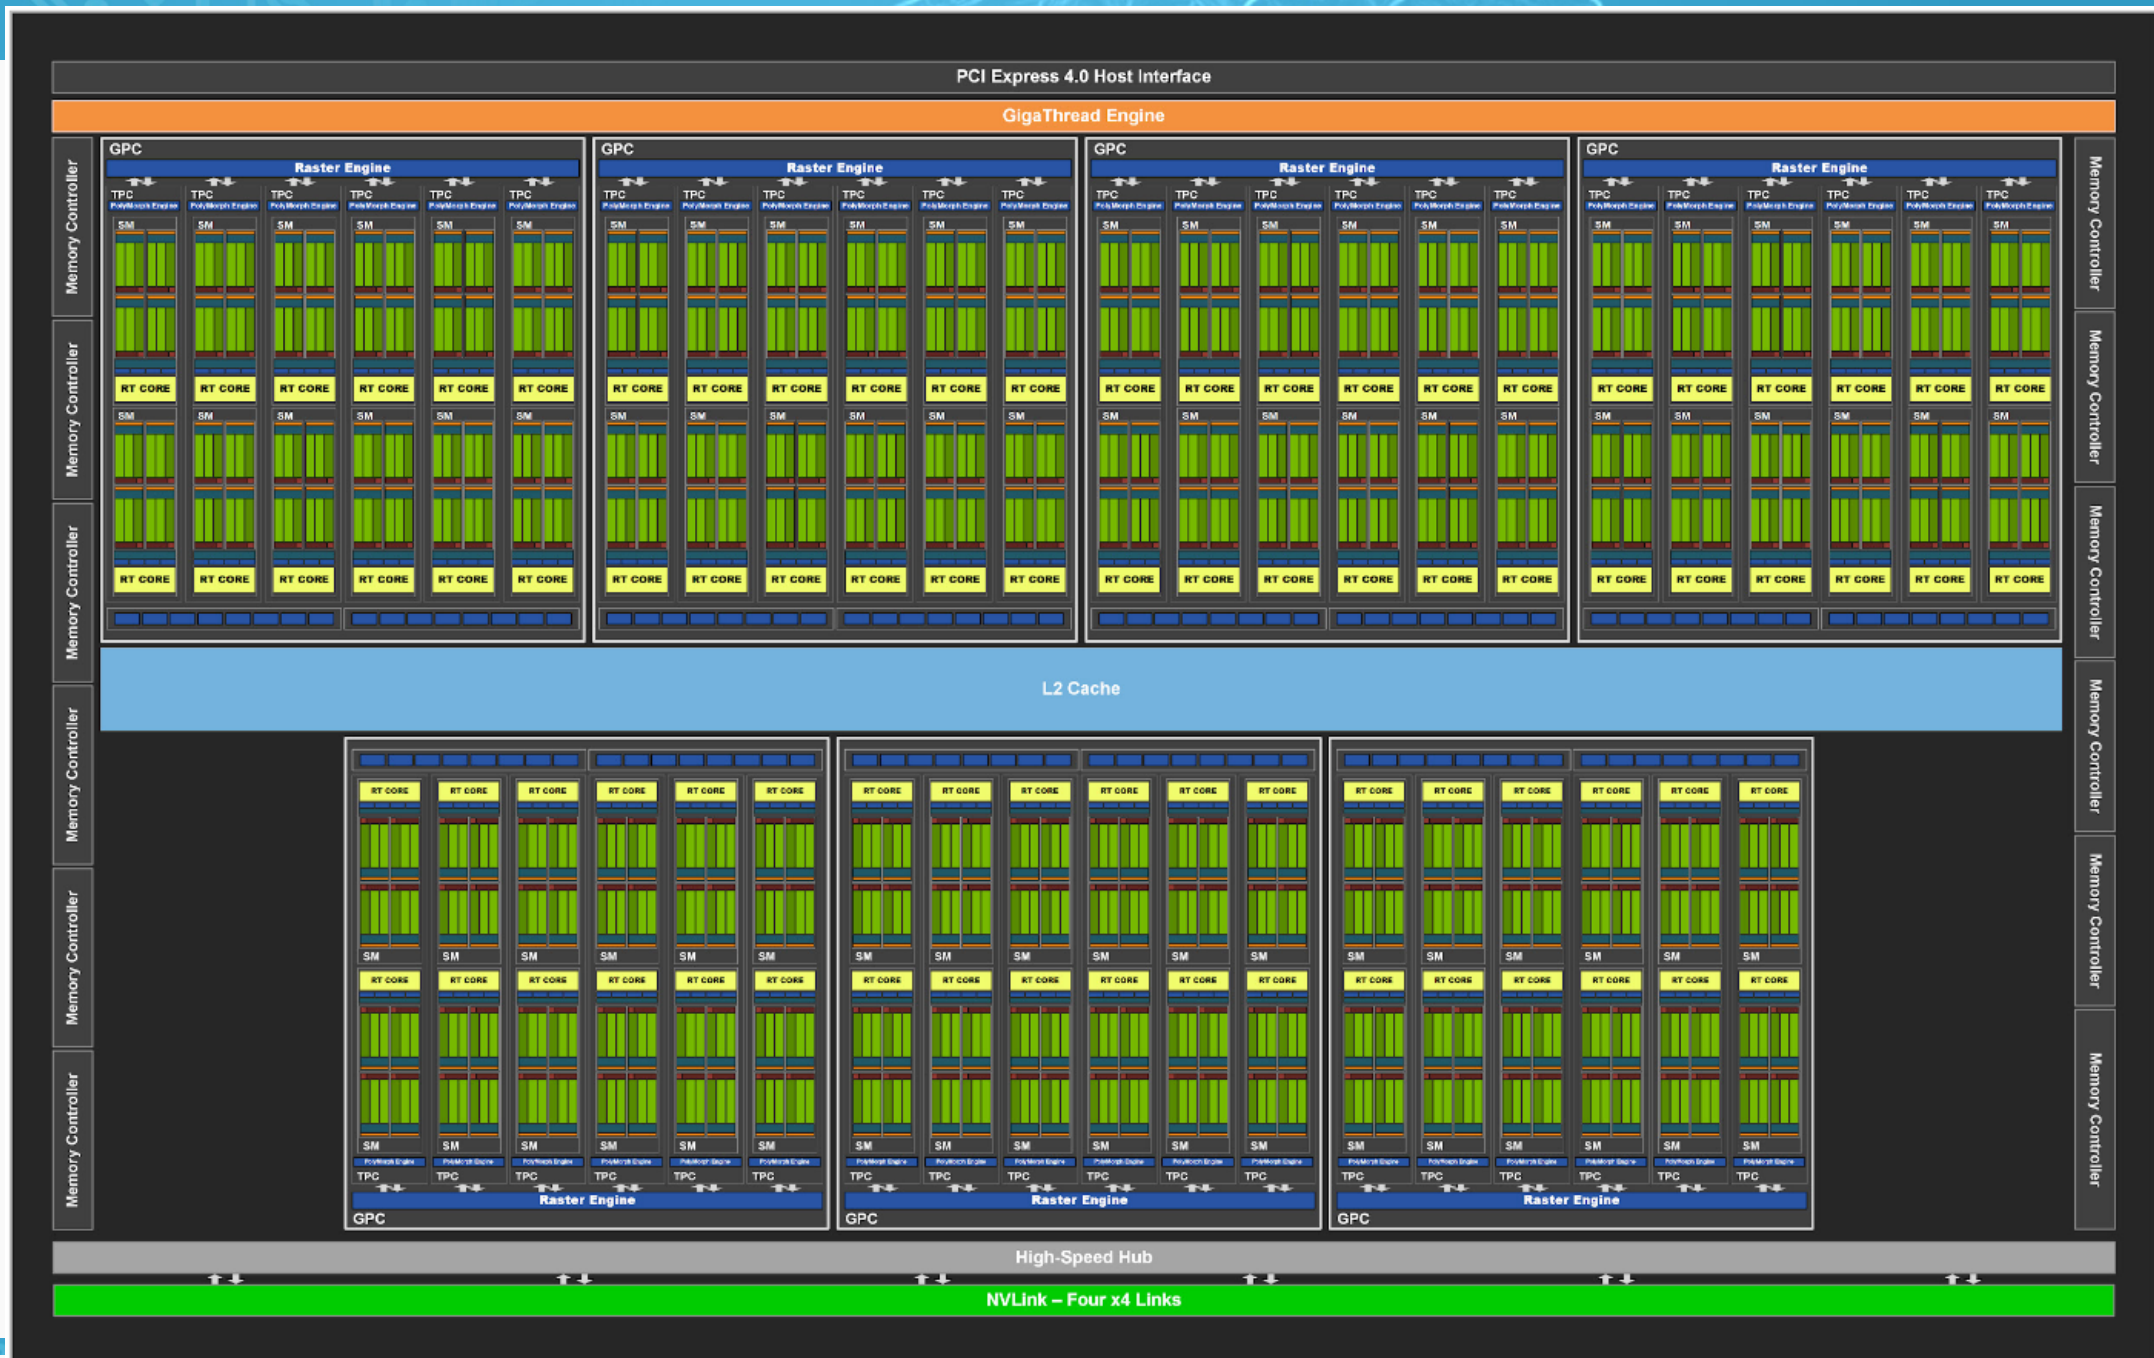
NVIDIA Ampere





NVIDIA Ampere v2







		GEFORCE RTX 3080 Ti	GEFORCE RTX 3080
Specificații GPU:	Nuclee NVIDIA CUDA®	10240	8704
	Frecvență Boost (GHz)	1.67	1.71
	Frecvență de bază (GHz)	1.37	1.44
Specificații memorie:	Configurație memorie standard	12 GB GDDR6X	10 GB GDDR6X
	Lățime interfață memorie	384 biți	320 biți
Tehnologii integrate:	Nuclee cu ray-tracing	Cea de-a doua generație	Cea de-a doua generație
	Nuclee Tensor	Cea de-a treia generație	Cea de-a treia generație
	Arhitectură NVIDIA	Ampere	Ampere



RTX 3080





NVIDIA A40

SPECIFICATIONS

GPU architecture	NVIDIA Ampere architecture
GPU memory	48 GB GDDR6 with ECC
Memory bandwidth	696 GB/s
Interconnect interface	NVIDIA® NVLink® 112.5 GB/s (bidirectional)³ PCIe Gen4 31.5 GB/s (bidirectional)
NVIDIA Ampere architecture- based CUDA Cores	10,752
NVIDIA second-generation RT Cores	84
NVIDIA third-generation Tensor Cores	336




CPU vs GPU?

- Cores
- Frequency
- Core complexity



Arhitectura system heterogen





Typical Program

- CPU alocă memorie pe GPU (în VRAM)
- CPU copiază date din RAM în VRAM
- CPU pornește **kernelul** pe GPU
- CPU copiază date din VRAM pe RAM





Open Standard for Parallel Programming of Heterogeneous Systems



OpenCLTM



OpenCL™



Desktop Creative Apps



Parallel Languages



Linear Algebra Libraries



Machine Learning Libraries and Frameworks



The industry's most pervasive, cross-vendor, open standard for low-level heterogeneous parallel programming

Molecular Modelling Libraries



Machine Learning Compilers



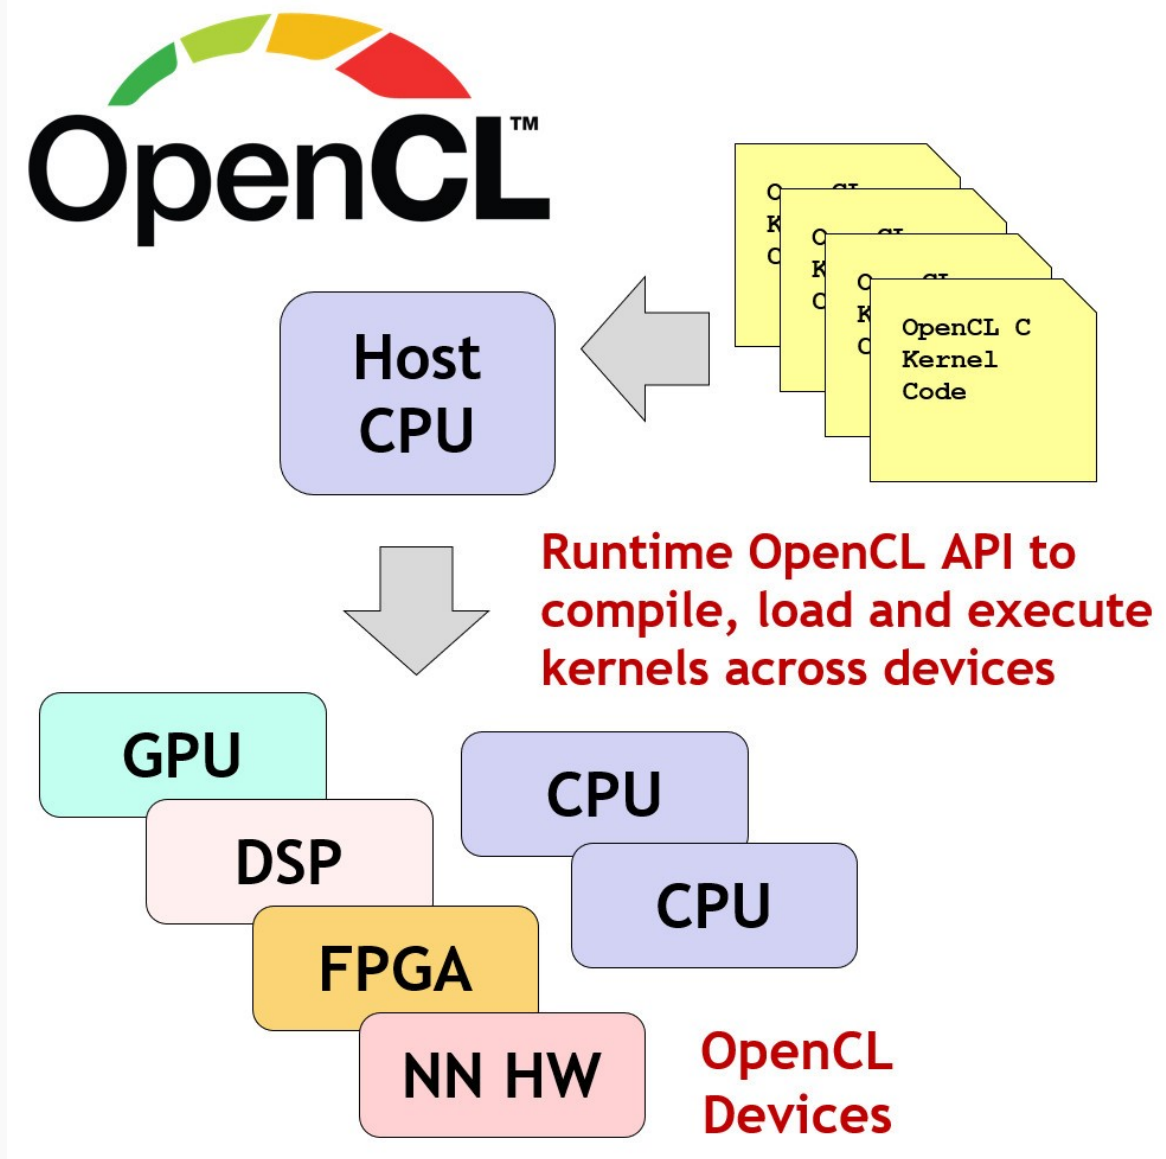
Vision, Imaging and Video Libraries



Math and Physics Libraries

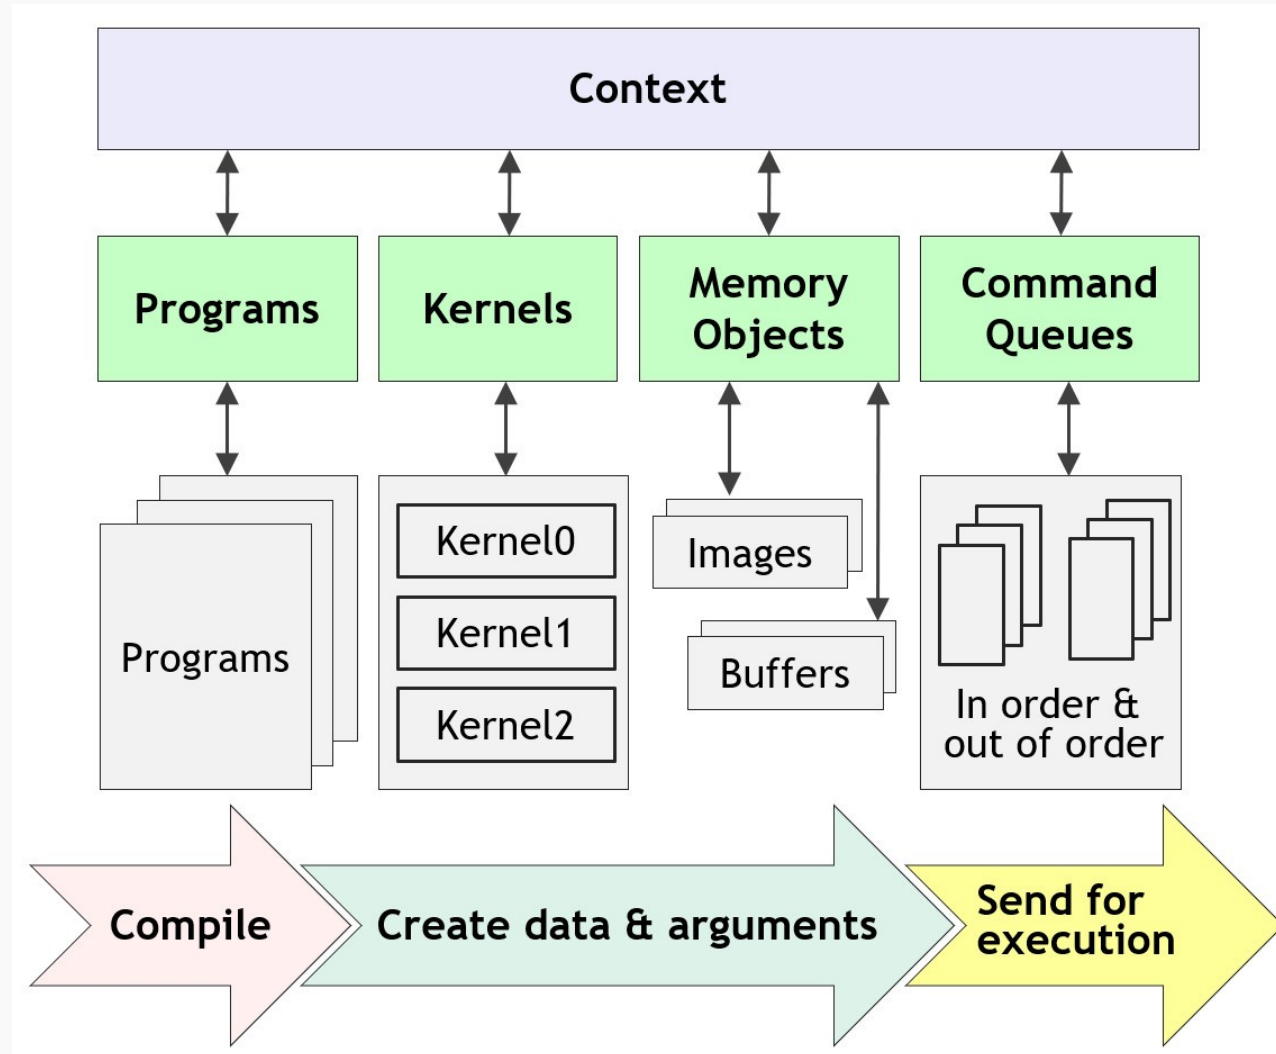


Accelerated Implementations





Sequence for Executing OpenCL Kernels





A complete sequence for executing an OpenCL program

- Query for available OpenCL platforms and devices
- Create a context for one or more OpenCL devices in a platform
- Create and build programs for OpenCL devices in the context
- Select kernels to execute from the programs
- Create memory objects for kernels to operate on
- Create command queues to execute commands on an OpenCL device
- *Enqueue* data transfer commands into the memory objects, if needed
- *Enqueue* kernels into the command queue for execution
- *Enqueue* commands to transfer data back to the host, if needed



Traditional Versus OpenCL Programming Using OpenCL C Kernels

C/C++ Programming

myapplication.c


// can also be a C++ file

hotspot 1

For (int i=0; i<N, i++) {...}

hotspot 2

For (int i=0; i<N, i++) {...}



CPU

OpenCL Programming

myapplication.c

// can also be a C++ file

clEnqueueNDRangeKernel(...);

clEnqueueNDRangeKernel(...);




device_code.cl

// OpenCL C - a C99 dialect

__kernel void k1() {
... //statements from hotspot 1 loop
}

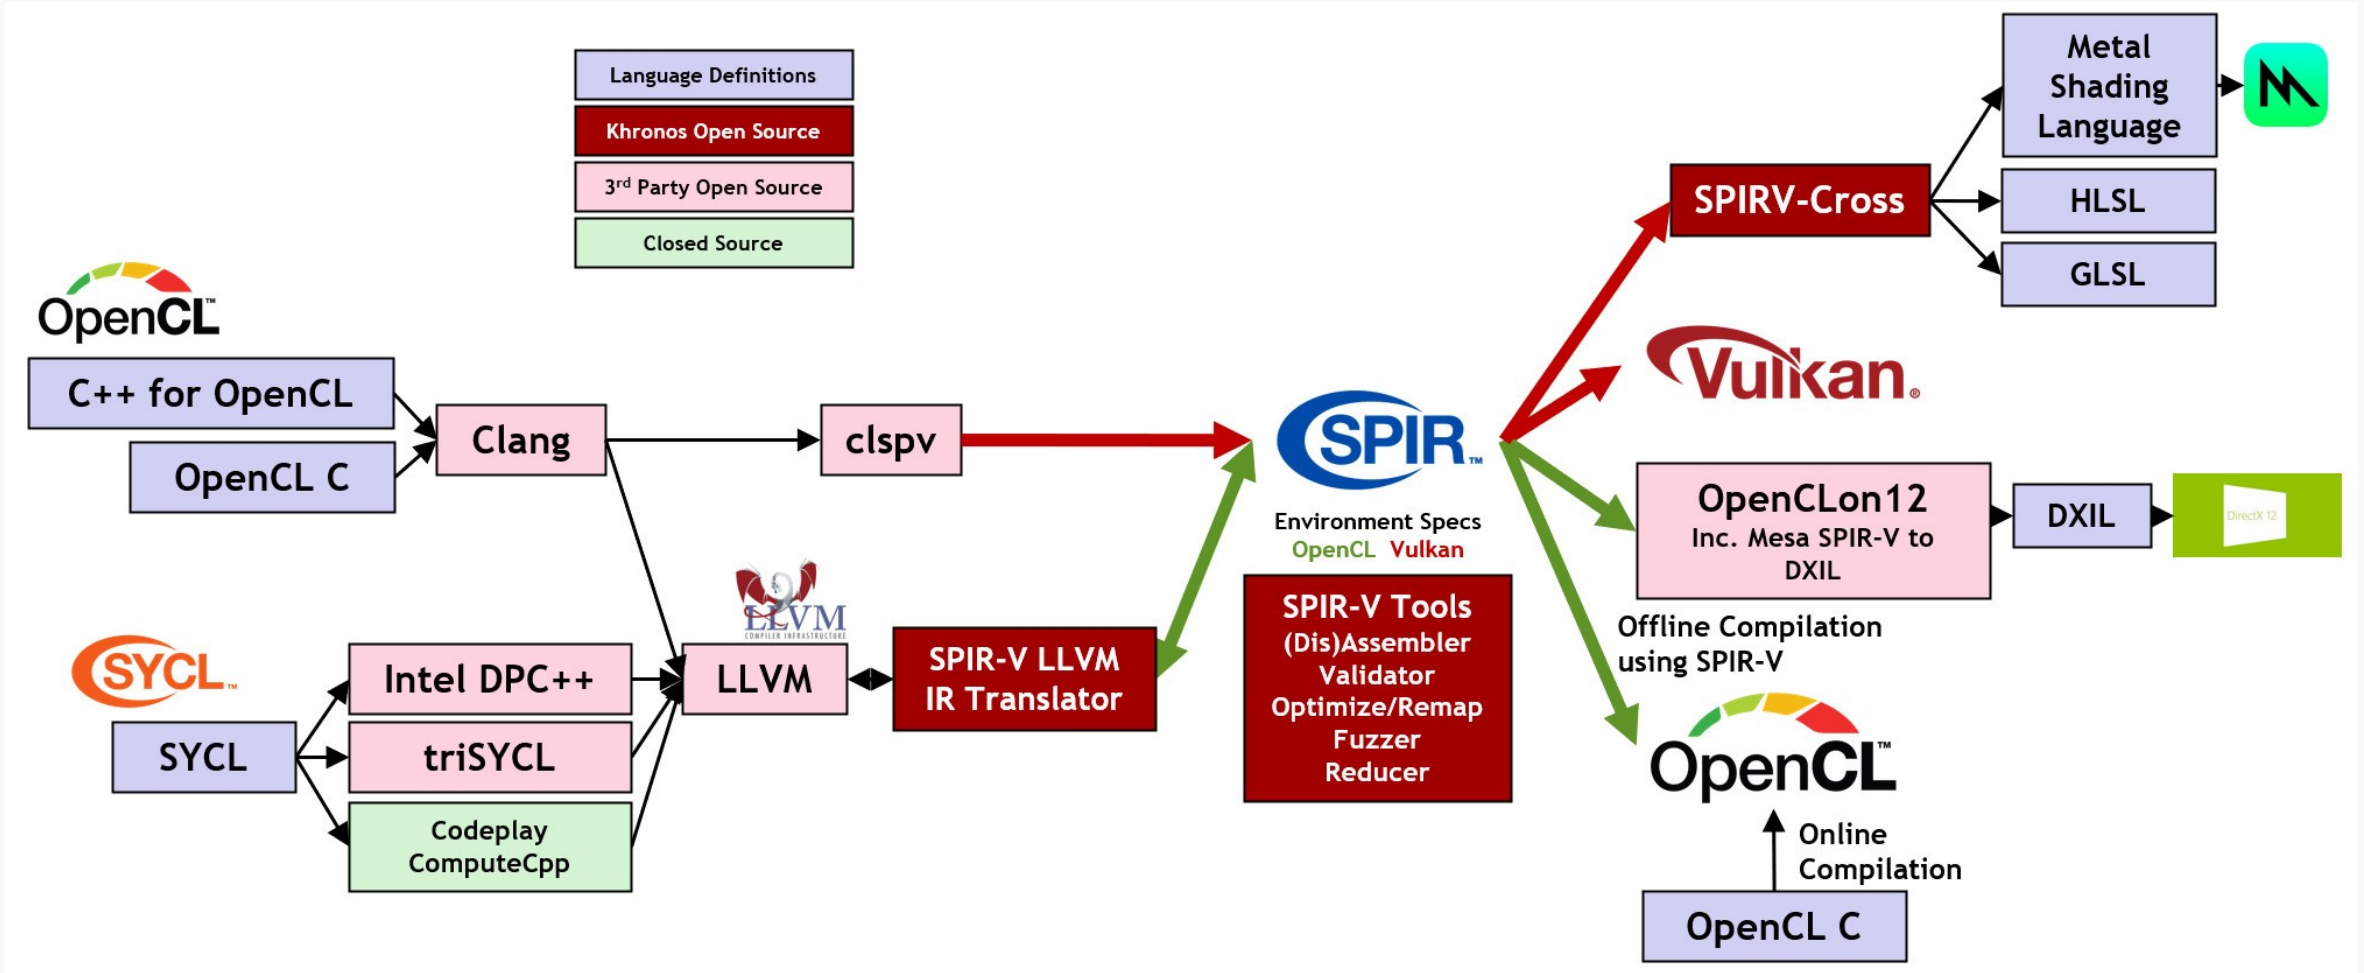
__kernel void k2() {
... //statements from hotspot 2 loop
}



Accelerator

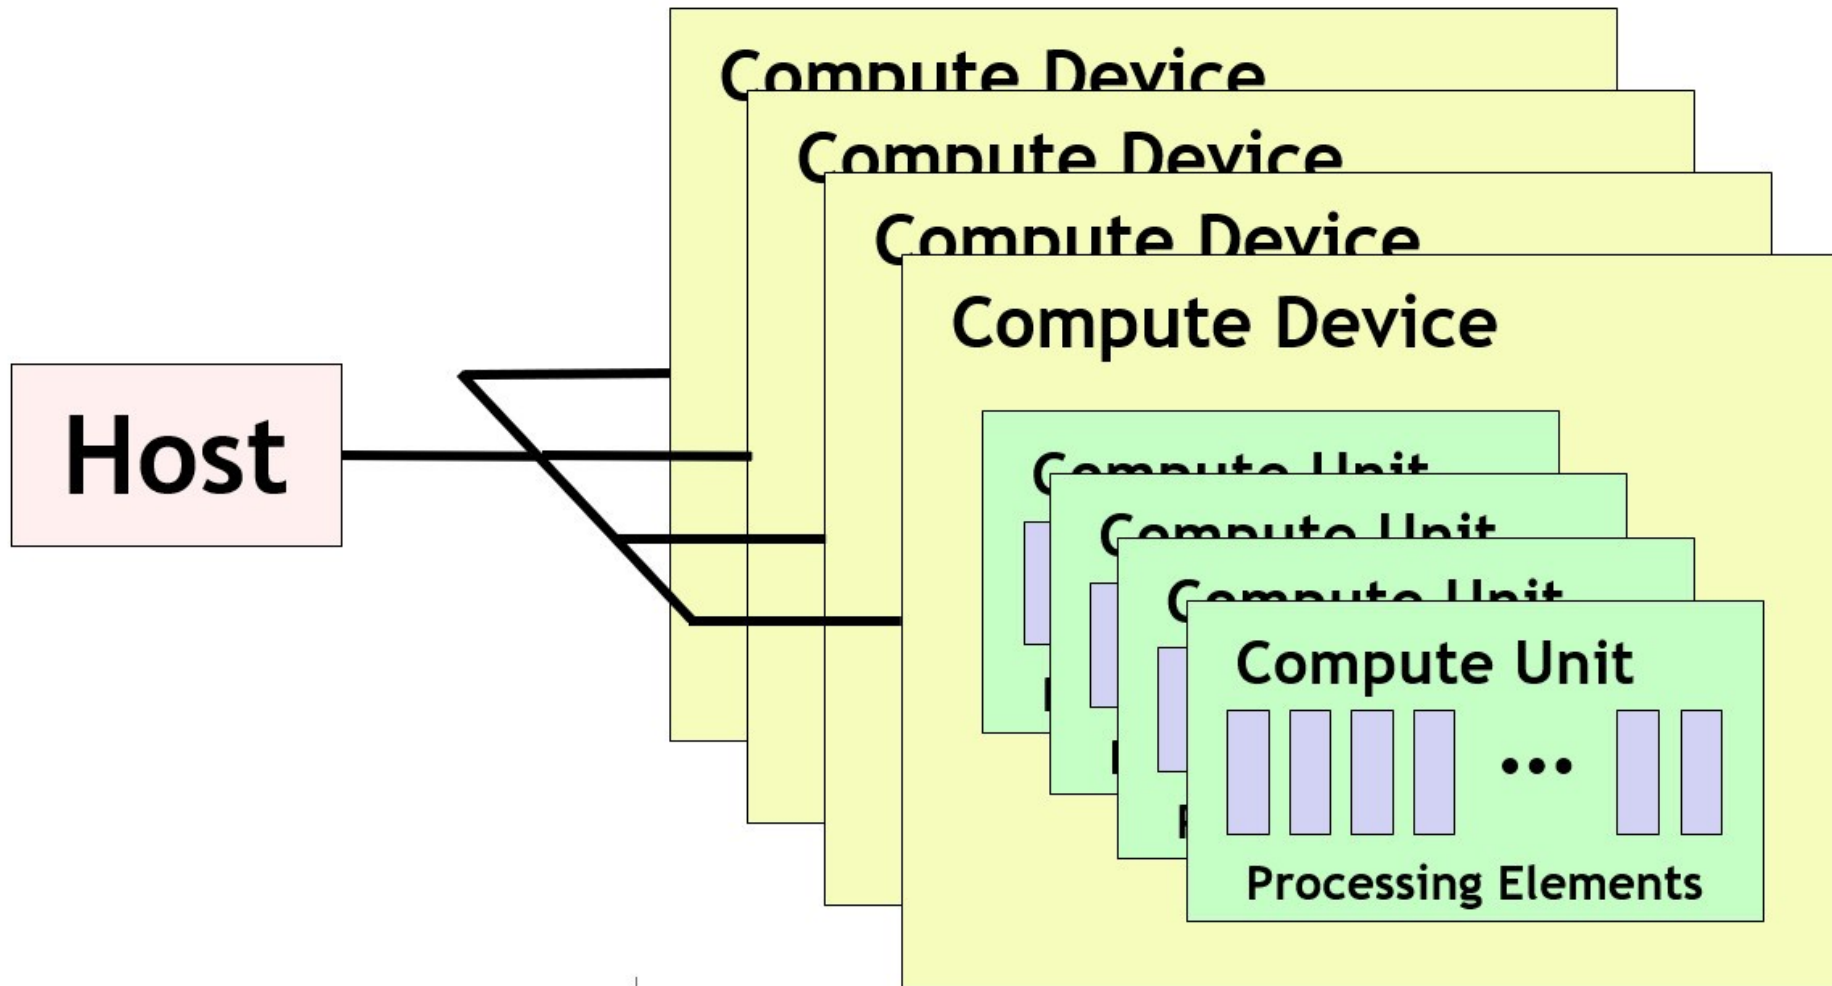


OpenCL Language Ecosystem Enabled With SPIR-V



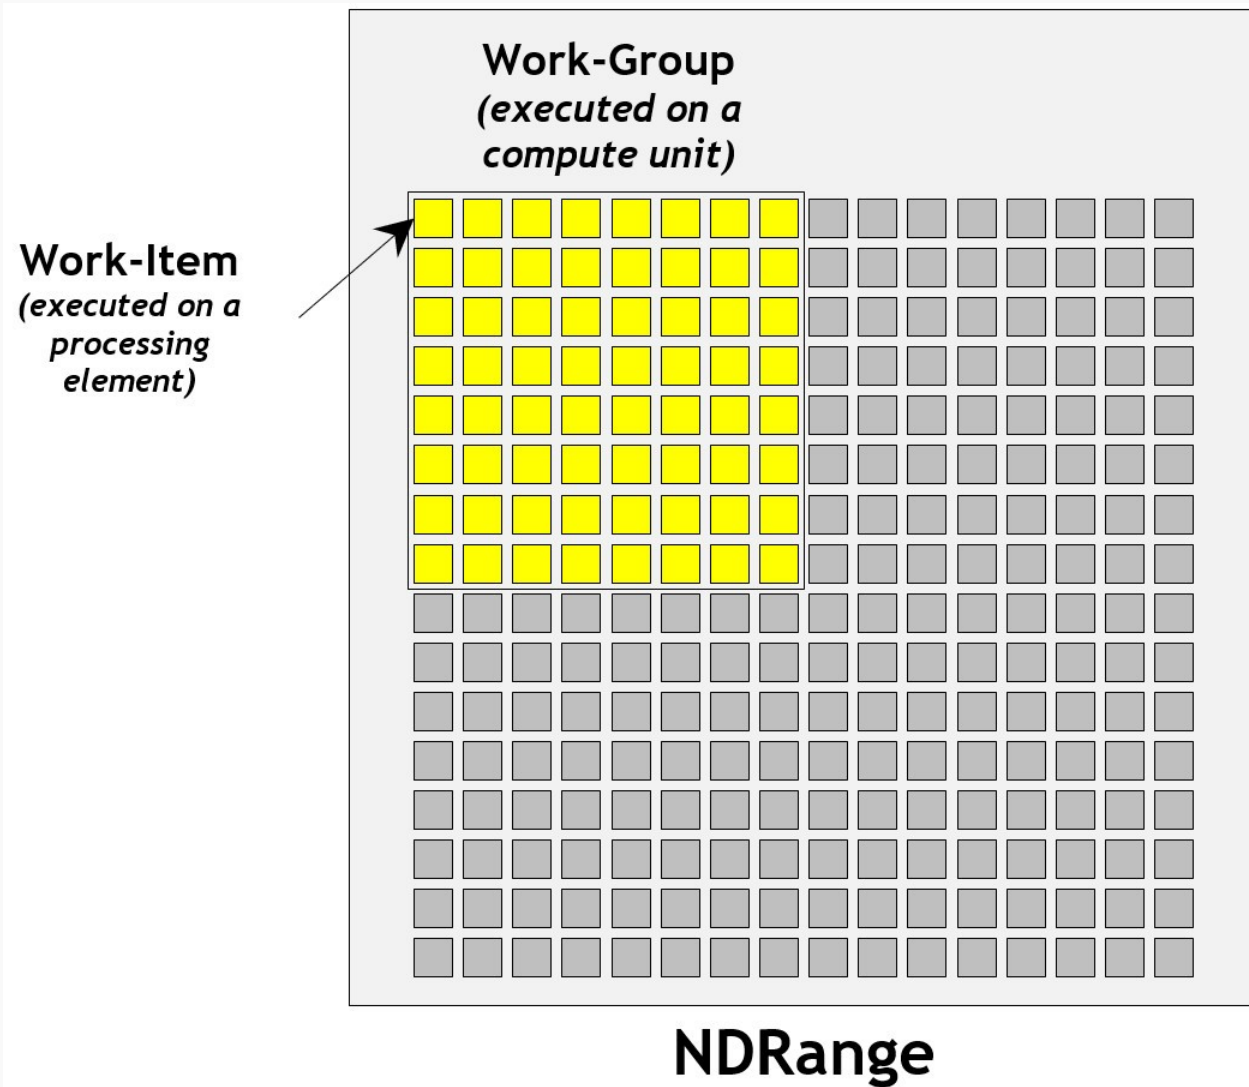


OpenCL Platform Model





A 2D Image as an Example NDRange





OpenCL Memory Model

